

```
In [32]: # !pip install tensorflow as tf
```

```
In [33]: # @Secure_Voice_Channel
def generic_vns_function(input_dim, number_dense_layers, classes, units):

    # adjusted to effectively import the libraries
    from keras.models import Sequential
    from keras.layers import Dense
    model = Sequential()

    # adjusted to leverage the correct functions
    for i in range(number_dense_layers):
        model.add(Dense(units=units, input_dim=input_dim,
                        kernel_initializer='normal',
                        activation='relu'))

    model.add(Dense(classes, kernel_initializer='normal',
                    activation='softmax'))
    model.compile(loss='categorical_crossentropy', optimizer='adam',
                  metrics=['accuracy'])
    return model

# adjusted to ingest varying number of epochs, batch sizes and units
def train_model(number_dense_layers, X_train, y_train, X_test, y_test,
                epochs, batch_size, units):
    # adjusted to leverage the generic_vns_function function
    model = generic_vns_function(X_train.shape[1], number_dense_layers,
                                y_train.shape[1], units)
    model.fit(X_train, y_train, validation_data=(X_test, y_test),
              epochs=epochs, batch_size=batch_size, verbose=2)
    scores = model.evaluate(X_test, y_test, verbose=2)
    print("Baseline Error: %.2f%%" % (100-scores[1]*100))
    return model
```

Computer vision

```
In [26]: # load data
from keras.datasets import mnist

(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [27]: # generic_vns_function(input_dim=num_pixels, number_dense_layers=1,
# @ classes=classes, units=1000)
```

```
In [28]: # reshaping the dataset from 28x28 to 784
num_pixels = X_test.shape[1]*X_test.shape[2] # 28x28 pixels = 784
X_test = X_test.reshape(X_test.shape[0], num_pixels).astype('float32')
X_train = X_train.reshape(X_train.shape[0], num_pixels).astype('float32')
```

```
In [29]: # normalize inputs
X_train = X_train / 255
X_test = X_test / 255
```

```
In [30]: # one hot encode outputs
from tensorflow.keras.utils import to_categorical
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
classes = y_test.shape[1] # Number of possible classes, = 10.
```

```
In [81]: train_model(number_dense_layers=1, X_train=X_train, y_train=y_train,
X_test=X_test, y_test=y_test, epochs=10, batch_size=200,
units=1000)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

- 3s - loss: 0.2571 - accuracy: 0.9264 - val_loss: 0.1220 - val_accuracy: 0.9632

Epoch 2/10

- 3s - loss: 0.1006 - accuracy: 0.9710 - val_loss: 0.0889 - val_accuracy: 0.9729

Epoch 3/10

- 3s - loss: 0.0642 - accuracy: 0.9812 - val_loss: 0.0764 - val_accuracy: 0.9767

Epoch 4/10

- 3s - loss: 0.0429 - accuracy: 0.9882 - val_loss: 0.0645 - val_accuracy: 0.9795

Epoch 5/10

- 3s - loss: 0.0297 - accuracy: 0.9920 - val_loss: 0.0680 - val_accuracy: 0.9791

Epoch 6/10

- 3s - loss: 0.0228 - accuracy: 0.9937 - val_loss: 0.0636 - val_accuracy: 0.9791

Epoch 7/10

- 3s - loss: 0.0155 - accuracy: 0.9963 - val_loss: 0.0584 - val_accuracy: 0.9813

Epoch 8/10

- 3s - loss: 0.0113 - accuracy: 0.9974 - val_loss: 0.0696 - val_accuracy: 0.9789

Epoch 9/10

- 3s - loss: 0.0093 - accuracy: 0.9979 - val_loss: 0.0613 - val_accuracy: 0.9821

Epoch 10/10

- 3s - loss: 0.0064 - accuracy: 0.9990 - val_loss: 0.0584 - val_accuracy: 0.9826

Baseline Error: 1.74%

```
Out[81]: <keras.engine.sequential.Sequential at 0x7f859d5ecbd0>
```

Vision Observations from default model

- Each epoch loss decreases, and accuracy increases on the training dataset
- There seems to be a bit of warning sign around epoch 8: accuracy on the training set went up, but accuracy on the validation set went down.

```
In [97]: # adding an additional dense layer
train_model(number_dense_layers=2, X_train= X_train, y_train=y_train,
            X_test=X_test, y_test= y_test, epochs=10, batch_size=200,
            units=1000)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

- 6s - loss: 0.1973 - accuracy: 0.9403 - val_loss: 0.1067 - val_accuracy: 0.9649

Epoch 2/10

- 6s - loss: 0.0693 - accuracy: 0.9786 - val_loss: 0.0746 - val_accuracy: 0.9769

Epoch 3/10

- 6s - loss: 0.0436 - accuracy: 0.9865 - val_loss: 0.0637 - val_accuracy: 0.9789

Epoch 4/10

- 6s - loss: 0.0290 - accuracy: 0.9905 - val_loss: 0.0696 - val_accuracy: 0.9800

Epoch 5/10

- 6s - loss: 0.0206 - accuracy: 0.9938 - val_loss: 0.0805 - val_accuracy: 0.9767

Epoch 6/10

- 6s - loss: 0.0189 - accuracy: 0.9938 - val_loss: 0.0768 - val_accuracy: 0.9782

Epoch 7/10

- 6s - loss: 0.0151 - accuracy: 0.9951 - val_loss: 0.0946 - val_accuracy: 0.9762

Epoch 8/10

- 6s - loss: 0.0143 - accuracy: 0.9953 - val_loss: 0.0883 - val_accuracy: 0.9783

Epoch 9/10

- 6s - loss: 0.0155 - accuracy: 0.9950 - val_loss: 0.0689 - val_accuracy: 0.9827

Epoch 10/10

- 6s - loss: 0.0113 - accuracy: 0.9965 - val_loss: 0.0858 - val_accuracy: 0.9794

Baseline Error: 2.06%

```
Out[97]: <keras.engine.sequential.Sequential at 0x7f7a4bc8ab50>
```

Conclusion: adding an additional layer did not increase accuracy for the vision model

In [98]: *# increasing epochs*

```
train_model(number_dense_layers=1, X_train= X_train, y_train=y_train,  
            X_test=X_test, y_test= y_test, epochs=20, batch_size=200,  
            units=1000)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

- 3s - loss: 0.2592 - accuracy: 0.9256 - val_loss: 0.1316 - val_accuracy: 0.9592

Epoch 2/20

- 3s - loss: 0.1007 - accuracy: 0.9703 - val_loss: 0.0898 - val_accuracy: 0.9733

Epoch 3/20

- 3s - loss: 0.0633 - accuracy: 0.9816 - val_loss: 0.0735 - val_accuracy: 0.9774

Epoch 4/20

- 3s - loss: 0.0433 - accuracy: 0.9877 - val_loss: 0.0618 - val_accuracy: 0.9793

Epoch 5/20

- 3s - loss: 0.0310 - accuracy: 0.9915 - val_loss: 0.0619 - val_accuracy: 0.9792

Epoch 6/20

- 3s - loss: 0.0220 - accuracy: 0.9944 - val_loss: 0.0654 - val_accuracy: 0.9794

Epoch 7/20

- 3s - loss: 0.0163 - accuracy: 0.9959 - val_loss: 0.0614 - val_accuracy: 0.9816

Epoch 8/20

- 3s - loss: 0.0121 - accuracy: 0.9974 - val_loss: 0.0588 - val_accuracy: 0.9820

Epoch 9/20

- 3s - loss: 0.0087 - accuracy: 0.9984 - val_loss: 0.0592 - val_accuracy: 0.9807

Epoch 10/20

- 3s - loss: 0.0058 - accuracy: 0.9991 - val_loss: 0.0577 - val_accuracy: 0.9826

Epoch 11/20

- 3s - loss: 0.0068 - accuracy: 0.9984 - val_loss: 0.0789 - val_accuracy: 0.9772

Epoch 12/20

- 3s - loss: 0.0081 - accuracy: 0.9979 - val_loss: 0.0662 - val_accuracy: 0.9815

Epoch 13/20

- 4s - loss: 0.0036 - accuracy: 0.9994 - val_loss: 0.0605 - val_accuracy: 0.9827

Epoch 14/20

- 3s - loss: 0.0023 - accuracy: 0.9998 - val_loss: 0.0643 - val_accuracy: 0.9820

Epoch 15/20

- 3s - loss: 0.0021 - accuracy: 0.9997 - val_loss: 0.0650 - val_accuracy: 0.9817

Epoch 16/20

- 3s - loss: 0.0019 - accuracy: 0.9997 - val_loss: 0.0812 - val_accuracy: 0.9800

Epoch 17/20

- 3s - loss: 0.0168 - accuracy: 0.9944 - val_loss: 0.0753 - val_accuracy: 0.9799

```
Epoch 18/20
- 3s - loss: 0.0067 - accuracy: 0.9981 - val_loss: 0.0638 - val_accu
racy: 0.9834
Epoch 19/20
- 3s - loss: 0.0018 - accuracy: 0.9997 - val_loss: 0.0599 - val_accu
racy: 0.9845
Epoch 20/20
- 3s - loss: 5.1167e-04 - accuracy: 1.0000 - val_loss: 0.0604 - val_accu
racy: 0.9850
Baseline Error: 1.50%
```

Out[98]: <keras.engine.sequential.Sequential at 0x7f7a4b3b7650>

Conclusion: increasing the number of epochs does increase accuracy for the vision model

```
In [99]: # decreasing batch size
```

```
train_model(number_dense_layers=1, X_train= X_train, y_train=y_train,  
            X_test=X_test, y_test= y_test, epochs=20, batch_size=100,  
            units=1000)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

- 5s - loss: 0.2231 - accuracy: 0.9347 - val_loss: 0.1150 - val_accuracy: 0.9665

Epoch 2/20

- 5s - loss: 0.0854 - accuracy: 0.9747 - val_loss: 0.0791 - val_accuracy: 0.9753

Epoch 3/20

- 5s - loss: 0.0536 - accuracy: 0.9835 - val_loss: 0.0673 - val_accuracy: 0.9795

Epoch 4/20

- 5s - loss: 0.0344 - accuracy: 0.9898 - val_loss: 0.0714 - val_accuracy: 0.9776

Epoch 5/20

- 5s - loss: 0.0247 - accuracy: 0.9930 - val_loss: 0.0616 - val_accuracy: 0.9800

Epoch 6/20

- 5s - loss: 0.0180 - accuracy: 0.9948 - val_loss: 0.0618 - val_accuracy: 0.9812

Epoch 7/20

- 5s - loss: 0.0137 - accuracy: 0.9959 - val_loss: 0.0679 - val_accuracy: 0.9807

Epoch 8/20

- 5s - loss: 0.0104 - accuracy: 0.9970 - val_loss: 0.0662 - val_accuracy: 0.9819

Epoch 9/20

- 5s - loss: 0.0105 - accuracy: 0.9969 - val_loss: 0.0617 - val_accuracy: 0.9825

Epoch 10/20

- 5s - loss: 0.0075 - accuracy: 0.9981 - val_loss: 0.0654 - val_accuracy: 0.9818

Epoch 11/20

- 5s - loss: 0.0123 - accuracy: 0.9961 - val_loss: 0.0671 - val_accuracy: 0.9819

Epoch 12/20

- 4s - loss: 0.0062 - accuracy: 0.9981 - val_loss: 0.0686 - val_accuracy: 0.9832

Epoch 13/20

- 5s - loss: 0.0023 - accuracy: 0.9995 - val_loss: 0.0662 - val_accuracy: 0.9847

Epoch 14/20

- 5s - loss: 0.0092 - accuracy: 0.9969 - val_loss: 0.0764 - val_accuracy: 0.9812

Epoch 15/20

- 5s - loss: 0.0069 - accuracy: 0.9980 - val_loss: 0.0706 - val_accuracy: 0.9831

Epoch 16/20

- 5s - loss: 0.0034 - accuracy: 0.9989 - val_loss: 0.0672 - val_accuracy: 0.9839

Epoch 17/20

- 5s - loss: 0.0030 - accuracy: 0.9991 - val_loss: 0.0793 - val_accuracy: 0.9822

```
Epoch 18/20
- 5s - loss: 0.0073 - accuracy: 0.9976 - val_loss: 0.0853 - val_accu
racy: 0.9814
Epoch 19/20
- 5s - loss: 0.0046 - accuracy: 0.9984 - val_loss: 0.0943 - val_accu
racy: 0.9813
Epoch 20/20
- 5s - loss: 0.0043 - accuracy: 0.9987 - val_loss: 0.0773 - val_accu
racy: 0.9844
Baseline Error: 1.56%
```

Out[99]: <keras.engine.sequential.Sequential at 0x7f7a4b8d4f90>

Conclusion: decreasing batch size did not increase accuracy for the vision model

```
In [31]: # increasing epochs
```

```
train_model(number_dense_layers=1, X_train= X_train, y_train=y_train,  
            X_test=X_test, y_test= y_test, epochs=20, batch_size=200,  
            units=2000)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

- 6s - loss: 0.2255 - accuracy: 0.9326 - val_loss: 0.1198 - val_accuracy: 0.9647

Epoch 2/20

- 6s - loss: 0.0833 - accuracy: 0.9757 - val_loss: 0.0786 - val_accuracy: 0.9756

Epoch 3/20

- 7s - loss: 0.0521 - accuracy: 0.9845 - val_loss: 0.0726 - val_accuracy: 0.9765

Epoch 4/20

- 6s - loss: 0.0334 - accuracy: 0.9901 - val_loss: 0.0663 - val_accuracy: 0.9792

Epoch 5/20

- 6s - loss: 0.0211 - accuracy: 0.9944 - val_loss: 0.0709 - val_accuracy: 0.9778

Epoch 6/20

- 7s - loss: 0.0141 - accuracy: 0.9965 - val_loss: 0.0624 - val_accuracy: 0.9811

Epoch 7/20

- 7s - loss: 0.0104 - accuracy: 0.9976 - val_loss: 0.0602 - val_accuracy: 0.9826

Epoch 8/20

- 7s - loss: 0.0064 - accuracy: 0.9987 - val_loss: 0.0666 - val_accuracy: 0.9806

Epoch 9/20

- 7s - loss: 0.0082 - accuracy: 0.9980 - val_loss: 0.0661 - val_accuracy: 0.9818

Epoch 10/20

- 7s - loss: 0.0089 - accuracy: 0.9974 - val_loss: 0.0684 - val_accuracy: 0.9810

Epoch 11/20

- 7s - loss: 0.0080 - accuracy: 0.9978 - val_loss: 0.0745 - val_accuracy: 0.9809

Epoch 12/20

- 7s - loss: 0.0116 - accuracy: 0.9960 - val_loss: 0.0759 - val_accuracy: 0.9793

Epoch 13/20

- 7s - loss: 0.0068 - accuracy: 0.9980 - val_loss: 0.0677 - val_accuracy: 0.9825

Epoch 14/20

- 7s - loss: 0.0038 - accuracy: 0.9991 - val_loss: 0.0712 - val_accuracy: 0.9817

Epoch 15/20

- 7s - loss: 0.0023 - accuracy: 0.9995 - val_loss: 0.0657 - val_accuracy: 0.9825

Epoch 16/20

- 7s - loss: 0.0031 - accuracy: 0.9991 - val_loss: 0.0774 - val_accuracy: 0.9828

Epoch 17/20

- 7s - loss: 0.0044 - accuracy: 0.9988 - val_loss: 0.0859 - val_accuracy: 0.9782


```
Epoch 18/20
- 7s - loss: 0.0146 - accuracy: 0.9954 - val_loss: 0.0807 - val_accuracy: 0.9806
Epoch 19/20
- 6s - loss: 0.0069 - accuracy: 0.9977 - val_loss: 0.0810 - val_accuracy: 0.9822
Epoch 20/20
- 6s - loss: 0.0027 - accuracy: 0.9993 - val_loss: 0.0702 - val_accuracy: 0.9836
Baseline Error: 1.64%
```

```
Out[31]: <keras.engine.sequential.Sequential at 0x7f980f2a1410>
```

Conclusion:

- the best model turned to be a model with one dense layer and 20 epochs
- there doesn't seem to be an intuitive reason why neither increasing batch size or units can't improve the result. Setting hyperparameters so far seems more than an art than a science.

Speech recognition

```
In [7]: pip install python_speech_features
```

```
Requirement already satisfied: python_speech_features in /opt/anaconda3/envs/deep37/lib/python3.7/site-packages (0.6)
Note: you may need to restart the kernel to use updated packages.
```

```
In [13]: pip install pandas
```

```
Collecting pandas
  Downloading pandas-1.1.2-cp37-cp37m-macosx_10_9_x86_64.whl (10.4 MB)
    |████████████████████████████████████████| 10.4 MB 2.4 MB/s eta 0:00:01
Requirement already satisfied: python-dateutil>=2.7.3 in /opt/anaconda3/envs/deep37/lib/python3.7/site-packages (from pandas) (2.8.1)
Collecting pytz>=2017.2
  Using cached pytz-2020.1-py2.py3-none-any.whl (510 kB)
Requirement already satisfied: numpy>=1.15.4 in /opt/anaconda3/envs/deep37/lib/python3.7/site-packages (from pandas) (1.18.1)
Requirement already satisfied: six>=1.5 in /opt/anaconda3/envs/deep37/lib/python3.7/site-packages (from python-dateutil>=2.7.3->pandas) (1.14.0)
Installing collected packages: pytz, pandas
Successfully installed pandas-1.1.2 pytz-2020.1
Note: you may need to restart the kernel to use updated packages.
```

```
In [ ]: #import python_speech_features
# mfcc = python_speech_features.mfcc(file, rate)
```

```
In [11]: # import spoken digit dataset (get_speech_dataset)
from db_utils import get_speech_dataset
(X_train, y_train), (X_test, y_test), (X_val, y_val) = get_speech_dataset()
# (X_train, y_train), (X_test, y_test) = get_speech_dataset()
```

```
In [12]: # flatten the data
num_pixels = X_test.shape[1]*X_test.shape[2]
X_test = X_test.reshape(X_test.shape[0], num_pixels).astype('float32')
X_train = X_train.reshape(X_train.shape[0], num_pixels).astype('float32')
```

```
In [13]: # standardize the inputs
import numpy as np
mean = np.mean(X_train)
std = np.std(X_train)
X_train = (X_train-std)/mean
X_test = (X_test-std)/mean
```

```
In [14]: # one hot encode the outputs to avoid queueing due to proximity of values
from tensorflow.keras.utils import to_categorical
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
classes = y_test.shape[1] # Number of possible classes, = 10.
```

```
In [86]: train_model(number_dense_layers=1, X_train= X_train, y_train=y_train,  
                  X_test=X_test, y_test= y_test, epochs=10, batch_size=200,  
                  units=1000)
```

Train on 18620 samples, validate on 2552 samples

Epoch 1/10

- 3s - loss: 26.8836 - accuracy: 0.3362 - val_loss: 3.2326 - val_accuracy: 0.4189

Epoch 2/10

- 3s - loss: 2.0473 - accuracy: 0.4948 - val_loss: 1.9372 - val_accuracy: 0.5106

Epoch 3/10

- 3s - loss: 1.3198 - accuracy: 0.5996 - val_loss: 1.5285 - val_accuracy: 0.5803

Epoch 4/10

- 3s - loss: 1.0028 - accuracy: 0.6780 - val_loss: 1.2918 - val_accuracy: 0.6364

Epoch 5/10

- 3s - loss: 0.8033 - accuracy: 0.7401 - val_loss: 1.1430 - val_accuracy: 0.6642

Epoch 6/10

- 3s - loss: 0.6661 - accuracy: 0.7844 - val_loss: 1.0823 - val_accuracy: 0.6947

Epoch 7/10

- 3s - loss: 0.5638 - accuracy: 0.8145 - val_loss: 1.0386 - val_accuracy: 0.7092

Epoch 8/10

- 3s - loss: 0.4798 - accuracy: 0.8439 - val_loss: 0.9570 - val_accuracy: 0.7386

Epoch 9/10

- 3s - loss: 0.4087 - accuracy: 0.8653 - val_loss: 0.9334 - val_accuracy: 0.7496

Epoch 10/10

- 3s - loss: 0.3387 - accuracy: 0.8905 - val_loss: 0.9283 - val_accuracy: 0.7645

Baseline Error: 23.55%

```
Out[86]: <keras.engine.sequential.Sequential at 0x7f859de676d0>
```

Speech - Conclusions on default model

- The model we developed generalizes between computer vision and speech recognition
- In contrast to computer vision, the baseline error remains 23.55%.
- It is reasonable to assume running more epochs could help as both train and validation error are still decreasing
- It also seems reasonable to assume that other model specs could further improve the model.

```
In [104]: # adding an additional dense layer
train_model(number_dense_layers=2, X_train= X_train, y_train=y_train,
            X_test=X_test, y_test= y_test, epochs=10, batch_size=200,
            units=1000)
```

Train on 18620 samples, validate on 2552 samples

Epoch 1/10

- 4s - loss: 22.9293 - accuracy: 0.3347 - val_loss: 1.9929 - val_accuracy: 0.4412

Epoch 2/10

- 4s - loss: 1.4581 - accuracy: 0.5432 - val_loss: 1.5184 - val_accuracy: 0.5517

Epoch 3/10

- 4s - loss: 1.0155 - accuracy: 0.6739 - val_loss: 1.3201 - val_accuracy: 0.6148

Epoch 4/10

- 4s - loss: 0.7683 - accuracy: 0.7487 - val_loss: 1.1635 - val_accuracy: 0.6646

Epoch 5/10

- 4s - loss: 0.5859 - accuracy: 0.8044 - val_loss: 1.0080 - val_accuracy: 0.7104

Epoch 6/10

- 4s - loss: 0.4675 - accuracy: 0.8420 - val_loss: 1.0629 - val_accuracy: 0.7132

Epoch 7/10

- 4s - loss: 0.3709 - accuracy: 0.8748 - val_loss: 0.9579 - val_accuracy: 0.7382

Epoch 8/10

- 4s - loss: 0.2747 - accuracy: 0.9086 - val_loss: 0.9272 - val_accuracy: 0.7582

Epoch 9/10

- 4s - loss: 0.2252 - accuracy: 0.9231 - val_loss: 0.9172 - val_accuracy: 0.7680

Epoch 10/10

- 4s - loss: 0.1838 - accuracy: 0.9376 - val_loss: 0.9212 - val_accuracy: 0.7751

Baseline Error: 22.49%

```
Out[104]: <keras.engine.sequential.Sequential at 0x7f7a4d20bc10>
```

Conclusion: adding a layer improves the speech model

```
In [105]: # increasing epochs while using one layer
train_model(number_dense_layers=1, X_train= X_train, y_train=y_train,
            X_test=X_test, y_test= y_test, epochs=20, batch_size=200,
            units=1000)
```

Train on 18620 samples, validate on 2552 samples

Epoch 1/20

- 3s - loss: 28.4327 - accuracy: 0.3272 - val_loss: 3.2682 - val_accuracy: 0.4122

Epoch 2/20

- 3s - loss: 2.0190 - accuracy: 0.4762 - val_loss: 1.8003 - val_accuracy: 0.4765

Epoch 3/20

- 3s - loss: 1.3007 - accuracy: 0.5833 - val_loss: 1.4661 - val_accuracy: 0.5748

Epoch 4/20

- 3s - loss: 1.0065 - accuracy: 0.6689 - val_loss: 1.3271 - val_accuracy: 0.6148

Epoch 5/20

- 3s - loss: 0.8277 - accuracy: 0.7251 - val_loss: 1.2889 - val_accuracy: 0.6399

Epoch 6/20

- 3s - loss: 0.7078 - accuracy: 0.7653 - val_loss: 1.0619 - val_accuracy: 0.6842

Epoch 7/20

- 3s - loss: 0.5941 - accuracy: 0.8014 - val_loss: 1.0487 - val_accuracy: 0.6920

Epoch 8/20

- 3s - loss: 0.4959 - accuracy: 0.8363 - val_loss: 0.9693 - val_accuracy: 0.7257

Epoch 9/20

- 3s - loss: 0.4335 - accuracy: 0.8550 - val_loss: 0.9776 - val_accuracy: 0.7237

Epoch 10/20

- 3s - loss: 0.3765 - accuracy: 0.8758 - val_loss: 0.9247 - val_accuracy: 0.7363

Epoch 11/20

- 3s - loss: 0.3342 - accuracy: 0.8898 - val_loss: 0.9370 - val_accuracy: 0.7469

Epoch 12/20

- 3s - loss: 0.2957 - accuracy: 0.9043 - val_loss: 0.9256 - val_accuracy: 0.7598

Epoch 13/20

- 3s - loss: 0.2594 - accuracy: 0.9139 - val_loss: 0.9339 - val_accuracy: 0.7629

Epoch 14/20

- 3s - loss: 0.2288 - accuracy: 0.9248 - val_loss: 0.9318 - val_accuracy: 0.7598

Epoch 15/20

- 3s - loss: 0.2086 - accuracy: 0.9332 - val_loss: 0.9046 - val_accuracy: 0.7759

Epoch 16/20

- 3s - loss: 0.1777 - accuracy: 0.9430 - val_loss: 0.9177 - val_accuracy: 0.7806

Epoch 17/20

- 3s - loss: 0.1540 - accuracy: 0.9526 - val_loss: 0.9221 - val_accuracy: 0.7884

```
Epoch 18/20
- 3s - loss: 0.1391 - accuracy: 0.9573 - val_loss: 0.9266 - val_accu
racy: 0.7915
Epoch 19/20
- 3s - loss: 0.1323 - accuracy: 0.9595 - val_loss: 0.9683 - val_accu
racy: 0.7845
Epoch 20/20
- 3s - loss: 0.1216 - accuracy: 0.9633 - val_loss: 0.9257 - val_accu
racy: 0.7911
Baseline Error: 20.89%
```

```
Out[105]: <keras.engine.sequential.Sequential at 0x7f7a4eb1c710>
```

Conclusion: increasing epochs (without increasing the model) increases accuracy. The model also doesn't seem to be done learning, so we'll increase the number of epochs in the next model. We'll also combine more epochs with two layers.

```
In [7]: # two layers; increased epochs
train_model(number_dense_layers=2, X_train= X_train, y_train=y_train,
            X_test=X_test, y_test= y_test, epochs=30, batch_size=200,
            units=1000)
```

Using TensorFlow backend.

Train on 18620 samples, validate on 2552 samples

Epoch 1/30

- 4s - loss: 22.8555 - accuracy: 0.3086 - val_loss: 1.9293 - val_accuracy: 0.4306

Epoch 2/30

- 4s - loss: 1.4847 - accuracy: 0.5312 - val_loss: 1.5164 - val_accuracy: 0.5525

Epoch 3/30

- 4s - loss: 1.0594 - accuracy: 0.6601 - val_loss: 1.2104 - val_accuracy: 0.6309

Epoch 4/30

- 4s - loss: 0.7857 - accuracy: 0.7429 - val_loss: 1.0193 - val_accuracy: 0.6799

Epoch 5/30

- 4s - loss: 0.6033 - accuracy: 0.8061 - val_loss: 0.9331 - val_accuracy: 0.7261

Epoch 6/30

- 4s - loss: 0.4828 - accuracy: 0.8413 - val_loss: 0.8836 - val_accuracy: 0.7390

Epoch 7/30

- 4s - loss: 0.4173 - accuracy: 0.8633 - val_loss: 0.8749 - val_accuracy: 0.7520

Epoch 8/30

- 4s - loss: 0.3311 - accuracy: 0.8908 - val_loss: 0.8320 - val_accuracy: 0.7696

Epoch 9/30

- 4s - loss: 0.2485 - accuracy: 0.9173 - val_loss: 0.7892 - val_accuracy: 0.7958

Epoch 10/30

- 4s - loss: 0.2153 - accuracy: 0.9266 - val_loss: 0.8490 - val_accuracy: 0.7739

Epoch 11/30

- 4s - loss: 0.1782 - accuracy: 0.9404 - val_loss: 0.8145 - val_accuracy: 0.8021

Epoch 12/30

- 4s - loss: 0.1568 - accuracy: 0.9476 - val_loss: 0.7907 - val_accuracy: 0.8037

Epoch 13/30

- 4s - loss: 0.1139 - accuracy: 0.9628 - val_loss: 0.8878 - val_accuracy: 0.7998

Epoch 14/30

- 4s - loss: 0.1218 - accuracy: 0.9596 - val_loss: 0.8368 - val_accuracy: 0.8103

Epoch 15/30

- 4s - loss: 0.1151 - accuracy: 0.9606 - val_loss: 0.8734 - val_accuracy: 0.8096

Epoch 16/30

- 4s - loss: 0.0901 - accuracy: 0.9709 - val_loss: 0.8813 - val_accuracy: 0.8194

Epoch 17/30

```
- 4s - loss: 0.0700 - accuracy: 0.9769 - val_loss: 0.8639 - val_accu-  
racy: 0.8190  
Epoch 18/30  
- 4s - loss: 0.0633 - accuracy: 0.9789 - val_loss: 0.9201 - val_accu-  
racy: 0.8143  
Epoch 19/30  
- 4s - loss: 0.0847 - accuracy: 0.9713 - val_loss: 1.0225 - val_accu-  
racy: 0.8107  
Epoch 20/30  
- 4s - loss: 0.0532 - accuracy: 0.9832 - val_loss: 0.9768 - val_accu-  
racy: 0.8186  
Epoch 21/30  
- 4s - loss: 0.0450 - accuracy: 0.9857 - val_loss: 1.0767 - val_accu-  
racy: 0.8009  
Epoch 22/30  
- 4s - loss: 0.0497 - accuracy: 0.9840 - val_loss: 1.1091 - val_accu-  
racy: 0.8068  
Epoch 23/30  
- 4s - loss: 0.1143 - accuracy: 0.9627 - val_loss: 1.0087 - val_accu-  
racy: 0.8139  
Epoch 24/30  
- 4s - loss: 0.1522 - accuracy: 0.9505 - val_loss: 1.7775 - val_accu-  
racy: 0.7484  
Epoch 25/30  
- 4s - loss: 0.4334 - accuracy: 0.8849 - val_loss: 1.4541 - val_accu-  
racy: 0.7567  
Epoch 26/30  
- 4s - loss: 0.2606 - accuracy: 0.9229 - val_loss: 0.9987 - val_accu-  
racy: 0.8088  
Epoch 27/30  
- 4s - loss: 0.1932 - accuracy: 0.9393 - val_loss: 0.9455 - val_accu-  
racy: 0.8264  
Epoch 28/30  
- 4s - loss: 0.1248 - accuracy: 0.9584 - val_loss: 0.9811 - val_accu-  
racy: 0.8268  
Epoch 29/30  
- 4s - loss: 0.0709 - accuracy: 0.9759 - val_loss: 0.9585 - val_accu-  
racy: 0.8327  
Epoch 30/30  
- 4s - loss: 0.0266 - accuracy: 0.9905 - val_loss: 0.9036 - val_accu-  
racy: 0.8413  
Baseline Error: 15.87%
```

```
Out[7]: <keras.engine.sequential.Sequential at 0x7fa3b4b18250>
```

Conclusion: model error dropped significantly from 20+ to 16+ percent.
Validation continues to go down consistently. We'll increase the epochs.


```
In [15]: # adding more epochs as validation accuracy is still increasing
train_model(number_dense_layers=2, X_train= X_train, y_train=y_train,
            X_test=X_test, y_test= y_test, epochs=40, batch_size=200,
            units=1000)
```

Train on 18620 samples, validate on 2552 samples

Epoch 1/40

- 4s - loss: 28.7320 - accuracy: 0.3277 - val_loss: 2.0146 - val_accuracy: 0.4224

Epoch 2/40

- 4s - loss: 1.4722 - accuracy: 0.5302 - val_loss: 1.5860 - val_accuracy: 0.5278

Epoch 3/40

- 4s - loss: 1.0978 - accuracy: 0.6434 - val_loss: 1.2976 - val_accuracy: 0.6101

Epoch 4/40

- 4s - loss: 0.8248 - accuracy: 0.7290 - val_loss: 1.1451 - val_accuracy: 0.6681

Epoch 5/40

- 4s - loss: 0.6575 - accuracy: 0.7840 - val_loss: 1.0502 - val_accuracy: 0.6967

Epoch 6/40

- 4s - loss: 0.5269 - accuracy: 0.8267 - val_loss: 0.9702 - val_accuracy: 0.7183

Epoch 7/40

- 4s - loss: 0.4284 - accuracy: 0.8612 - val_loss: 0.9807 - val_accuracy: 0.7226

Epoch 8/40

- 4s - loss: 0.3497 - accuracy: 0.8838 - val_loss: 0.8753 - val_accuracy: 0.7712

Epoch 9/40

- 4s - loss: 0.2888 - accuracy: 0.9050 - val_loss: 0.8916 - val_accuracy: 0.7641

Epoch 10/40

- 4s - loss: 0.2471 - accuracy: 0.9156 - val_loss: 0.9434 - val_accuracy: 0.7641

Epoch 11/40

- 4s - loss: 0.2001 - accuracy: 0.9337 - val_loss: 0.8954 - val_accuracy: 0.7778

Epoch 12/40

- 4s - loss: 0.1903 - accuracy: 0.9348 - val_loss: 0.9532 - val_accuracy: 0.7853

Epoch 13/40

- 4s - loss: 0.1306 - accuracy: 0.9578 - val_loss: 0.9695 - val_accuracy: 0.7798

Epoch 14/40

- 4s - loss: 0.1105 - accuracy: 0.9649 - val_loss: 0.8898 - val_accuracy: 0.7970

Epoch 15/40

- 4s - loss: 0.1095 - accuracy: 0.9640 - val_loss: 0.9157 - val_accuracy: 0.8025

Epoch 16/40

- 4s - loss: 0.0869 - accuracy: 0.9718 - val_loss: 0.9532 - val_accuracy: 0.7955

Epoch 17/40

- 4s - loss: 0.0701 - accuracy: 0.9782 - val_loss: 0.9095 - val_accuracy: 0.8115

```
Epoch 18/40
- 4s - loss: 0.0469 - accuracy: 0.9867 - val_loss: 1.0180 - val_accuracy: 0.8002
Epoch 19/40
- 4s - loss: 0.0510 - accuracy: 0.9832 - val_loss: 0.9875 - val_accuracy: 0.8115
Epoch 20/40
- 4s - loss: 0.0440 - accuracy: 0.9873 - val_loss: 0.9434 - val_accuracy: 0.8217
Epoch 21/40
- 4s - loss: 0.0304 - accuracy: 0.9917 - val_loss: 1.0081 - val_accuracy: 0.8178
Epoch 22/40
- 4s - loss: 0.0689 - accuracy: 0.9774 - val_loss: 1.1835 - val_accuracy: 0.7861
Epoch 23/40
- 4s - loss: 0.1720 - accuracy: 0.9408 - val_loss: 1.1325 - val_accuracy: 0.7978
Epoch 24/40
- 4s - loss: 0.1920 - accuracy: 0.9396 - val_loss: 1.4916 - val_accuracy: 0.7551
Epoch 25/40
- 4s - loss: 0.3465 - accuracy: 0.9020 - val_loss: 1.2406 - val_accuracy: 0.7900
Epoch 26/40
- 4s - loss: 0.2753 - accuracy: 0.9190 - val_loss: 1.3598 - val_accuracy: 0.7813
Epoch 27/40
- 4s - loss: 0.2280 - accuracy: 0.9298 - val_loss: 1.1924 - val_accuracy: 0.8072
Epoch 28/40
- 4s - loss: 0.1195 - accuracy: 0.9577 - val_loss: 1.0743 - val_accuracy: 0.8135
Epoch 29/40
- 4s - loss: 0.1211 - accuracy: 0.9608 - val_loss: 1.2234 - val_accuracy: 0.8013
Epoch 30/40
- 4s - loss: 0.1247 - accuracy: 0.9599 - val_loss: 1.2023 - val_accuracy: 0.8135
Epoch 31/40
- 4s - loss: 0.0876 - accuracy: 0.9708 - val_loss: 1.1661 - val_accuracy: 0.8131
Epoch 32/40
- 4s - loss: 0.1036 - accuracy: 0.9663 - val_loss: 1.1898 - val_accuracy: 0.8217
Epoch 33/40
- 4s - loss: 0.0635 - accuracy: 0.9779 - val_loss: 1.1359 - val_accuracy: 0.8284
Epoch 34/40
- 4s - loss: 0.0534 - accuracy: 0.9825 - val_loss: 1.1293 - val_accuracy: 0.8225
Epoch 35/40
- 4s - loss: 0.0780 - accuracy: 0.9752 - val_loss: 1.3012 - val_accuracy: 0.8096
Epoch 36/40
- 4s - loss: 0.1010 - accuracy: 0.9688 - val_loss: 1.2635 - val_accuracy: 0.8197
```

```
Epoch 37/40
- 4s - loss: 0.0950 - accuracy: 0.9695 - val_loss: 1.1657 - val_accu-
racy: 0.8315
Epoch 38/40
- 4s - loss: 0.1196 - accuracy: 0.9641 - val_loss: 1.3705 - val_accu-
racy: 0.8033
Epoch 39/40
- 4s - loss: 0.2309 - accuracy: 0.9380 - val_loss: 1.5089 - val_accu-
racy: 0.7958
Epoch 40/40
- 4s - loss: 0.2138 - accuracy: 0.9433 - val_loss: 1.6747 - val_accu-
racy: 0.8013
Baseline Error: 19.87%
```

```
Out[15]: <keras.engine.sequential.Sequential at 0x7fa3b4b49310>
```

Conclusion: we observe signs of overfitting (decreasing validation errors) as of epoch 35. We'll stop the next model there, and explore batch size next.

```
In [18]: train_model(number_dense_layers=2, X_train= X_train, y_train=y_train,  
                    X_test=X_test, y_test= y_test, epochs=35, batch_size=100,  
                    units=1000)
```

Train on 18620 samples, validate on 2552 samples

Epoch 1/35

- 7s - loss: 14.3605 - accuracy: 0.3959 - val_loss: 1.5848 - val_accuracy: 0.5223

Epoch 2/35

- 6s - loss: 1.1489 - accuracy: 0.6321 - val_loss: 1.3096 - val_accuracy: 0.6105

Epoch 3/35

- 8s - loss: 0.7934 - accuracy: 0.7478 - val_loss: 1.0109 - val_accuracy: 0.6928

Epoch 4/35

- 7s - loss: 0.5902 - accuracy: 0.8072 - val_loss: 0.9551 - val_accuracy: 0.7198

Epoch 5/35

- 7s - loss: 0.4674 - accuracy: 0.8468 - val_loss: 0.8564 - val_accuracy: 0.7582

Epoch 6/35

- 7s - loss: 0.4090 - accuracy: 0.8645 - val_loss: 0.8032 - val_accuracy: 0.7782

Epoch 7/35

- 7s - loss: 0.3383 - accuracy: 0.8861 - val_loss: 0.8636 - val_accuracy: 0.7766

Epoch 8/35

- 6s - loss: 0.2861 - accuracy: 0.9038 - val_loss: 0.8317 - val_accuracy: 0.7947

Epoch 9/35

- 6s - loss: 0.2448 - accuracy: 0.9194 - val_loss: 0.9595 - val_accuracy: 0.7829

Epoch 10/35

- 6s - loss: 0.2617 - accuracy: 0.9142 - val_loss: 1.0185 - val_accuracy: 0.7731

Epoch 11/35

- 6s - loss: 0.2535 - accuracy: 0.9192 - val_loss: 1.0130 - val_accuracy: 0.7853

Epoch 12/35

- 6s - loss: 0.2712 - accuracy: 0.9129 - val_loss: 1.1181 - val_accuracy: 0.7798

Epoch 13/35

- 6s - loss: 0.2252 - accuracy: 0.9300 - val_loss: 0.9736 - val_accuracy: 0.8127

Epoch 14/35

- 6s - loss: 0.3075 - accuracy: 0.9088 - val_loss: 1.1095 - val_accuracy: 0.7915

Epoch 15/35

- 6s - loss: 0.2317 - accuracy: 0.9299 - val_loss: 1.0414 - val_accuracy: 0.8111

Epoch 16/35

- 6s - loss: 0.2065 - accuracy: 0.9332 - val_loss: 0.9959 - val_accuracy: 0.8170

Epoch 17/35

- 6s - loss: 0.1958 - accuracy: 0.9394 - val_loss: 1.2746 - val_accuracy: 0.7958

Epoch 18/35

```
- 6s - loss: 0.2684 - accuracy: 0.9222 - val_loss: 1.0934 - val_accuracy: 0.7962
Epoch 19/35
- 6s - loss: 0.2154 - accuracy: 0.9340 - val_loss: 1.0218 - val_accuracy: 0.8256
Epoch 20/35
- 6s - loss: 0.2307 - accuracy: 0.9322 - val_loss: 1.1134 - val_accuracy: 0.8174
Epoch 21/35
- 6s - loss: 0.2086 - accuracy: 0.9387 - val_loss: 1.1377 - val_accuracy: 0.8135
Epoch 22/35
- 6s - loss: 0.1719 - accuracy: 0.9493 - val_loss: 1.2245 - val_accuracy: 0.8154
Epoch 23/35
- 6s - loss: 0.2080 - accuracy: 0.9385 - val_loss: 1.1235 - val_accuracy: 0.8178
Epoch 24/35
- 7s - loss: 0.1416 - accuracy: 0.9555 - val_loss: 1.1198 - val_accuracy: 0.8382
Epoch 25/35
- 6s - loss: 0.1788 - accuracy: 0.9494 - val_loss: 0.9720 - val_accuracy: 0.8389
Epoch 26/35
- 6s - loss: 0.1339 - accuracy: 0.9579 - val_loss: 1.2324 - val_accuracy: 0.8150
Epoch 27/35
- 6s - loss: 0.1541 - accuracy: 0.9548 - val_loss: 1.1913 - val_accuracy: 0.8237
Epoch 28/35
- 6s - loss: 0.1675 - accuracy: 0.9519 - val_loss: 1.0991 - val_accuracy: 0.8370
Epoch 29/35
- 6s - loss: 0.1599 - accuracy: 0.9544 - val_loss: 1.2818 - val_accuracy: 0.8245
Epoch 30/35
- 6s - loss: 0.1317 - accuracy: 0.9598 - val_loss: 0.9771 - val_accuracy: 0.8417
Epoch 31/35
- 6s - loss: 0.1353 - accuracy: 0.9595 - val_loss: 1.2327 - val_accuracy: 0.8256
Epoch 32/35
- 6s - loss: 0.1602 - accuracy: 0.9565 - val_loss: 1.0702 - val_accuracy: 0.8487
Epoch 33/35
- 6s - loss: 0.1002 - accuracy: 0.9690 - val_loss: 1.2223 - val_accuracy: 0.8370
Epoch 34/35
- 6s - loss: 0.1457 - accuracy: 0.9579 - val_loss: 1.1952 - val_accuracy: 0.8484
Epoch 35/35
- 6s - loss: 0.1211 - accuracy: 0.9647 - val_loss: 1.1628 - val_accuracy: 0.8460
Baseline Error: 15.40%
```

Out[18]: <keras.engine.sequential.Sequential at 0x7f9b76b17ed0>

Conclusion: lowering batch size seems to improve the model mildly.

```
In [20]: train_model(number_dense_layers=2, X_train= X_train, y_train=y_train,  
                    X_test=X_test, y_test= y_test, epochs=35, batch_size=100,  
                    units=2000)
```

Train on 18620 samples, validate on 2552 samples

Epoch 1/35

- 18s - loss: 51.2224 - accuracy: 0.3471 - val_loss: 1.8722 - val_accuracy: 0.4165

Epoch 2/35

- 18s - loss: 1.4065 - accuracy: 0.5387 - val_loss: 1.5753 - val_accuracy: 0.5255

Epoch 3/35

- 18s - loss: 1.0712 - accuracy: 0.6488 - val_loss: 1.2885 - val_accuracy: 0.6164

Epoch 4/35

- 18s - loss: 0.8431 - accuracy: 0.7248 - val_loss: 1.2451 - val_accuracy: 0.6528

Epoch 5/35

- 18s - loss: 0.7023 - accuracy: 0.7705 - val_loss: 1.1592 - val_accuracy: 0.6814

Epoch 6/35

- 18s - loss: 0.5708 - accuracy: 0.8150 - val_loss: 1.1567 - val_accuracy: 0.7081

Epoch 7/35

- 18s - loss: 0.5881 - accuracy: 0.8125 - val_loss: 1.1787 - val_accuracy: 0.7116

Epoch 8/35

- 18s - loss: 0.5060 - accuracy: 0.8424 - val_loss: 1.1888 - val_accuracy: 0.7339

Epoch 9/35

- 18s - loss: 0.5177 - accuracy: 0.8416 - val_loss: 1.1389 - val_accuracy: 0.7394

Epoch 10/35

- 17s - loss: 0.4844 - accuracy: 0.8508 - val_loss: 1.0547 - val_accuracy: 0.7680

Epoch 11/35

- 18s - loss: 0.4275 - accuracy: 0.8704 - val_loss: 1.1926 - val_accuracy: 0.7653

Epoch 12/35

- 20s - loss: 0.4038 - accuracy: 0.8795 - val_loss: 1.1648 - val_accuracy: 0.7621

Epoch 13/35

- 18s - loss: 0.4808 - accuracy: 0.8622 - val_loss: 1.1720 - val_accuracy: 0.7594

Epoch 14/35

- 19s - loss: 0.4381 - accuracy: 0.8785 - val_loss: 1.3413 - val_accuracy: 0.7665

Epoch 15/35

- 18s - loss: 0.3996 - accuracy: 0.8861 - val_loss: 1.2496 - val_accuracy: 0.7786

Epoch 16/35

- 18s - loss: 0.3506 - accuracy: 0.9001 - val_loss: 1.2700 - val_accuracy: 0.7829

Epoch 17/35

- 18s - loss: 0.3608 - accuracy: 0.9016 - val_loss: 1.1891 - val_accuracy: 0.7810

Epoch 18/35

```
- 18s - loss: 0.4199 - accuracy: 0.8914 - val_loss: 1.3686 - val_accu-  
racy: 0.7708  
Epoch 19/35  
- 18s - loss: 0.4007 - accuracy: 0.8959 - val_loss: 1.3264 - val_accu-  
racy: 0.7782  
Epoch 20/35  
- 18s - loss: 0.3722 - accuracy: 0.9017 - val_loss: 1.5191 - val_accu-  
racy: 0.7606  
Epoch 21/35  
- 19s - loss: 0.3267 - accuracy: 0.9122 - val_loss: 1.2879 - val_accu-  
racy: 0.7915  
Epoch 22/35  
- 18s - loss: 0.3093 - accuracy: 0.9150 - val_loss: 1.2363 - val_accu-  
racy: 0.8025  
Epoch 23/35  
- 18s - loss: 0.2410 - accuracy: 0.9297 - val_loss: 1.1787 - val_accu-  
racy: 0.8056  
Epoch 24/35  
- 18s - loss: 0.2377 - accuracy: 0.9308 - val_loss: 1.2409 - val_accu-  
racy: 0.8072  
Epoch 25/35  
- 18s - loss: 0.2595 - accuracy: 0.9311 - val_loss: 1.3103 - val_accu-  
racy: 0.8131  
Epoch 26/35  
- 19s - loss: 0.2133 - accuracy: 0.9369 - val_loss: 1.2300 - val_accu-  
racy: 0.8068  
Epoch 27/35  
- 19s - loss: 0.2501 - accuracy: 0.9302 - val_loss: 1.5590 - val_accu-  
racy: 0.7896  
Epoch 28/35  
- 21s - loss: 0.2499 - accuracy: 0.9338 - val_loss: 1.3390 - val_accu-  
racy: 0.8096  
Epoch 29/35  
- 18s - loss: 0.1947 - accuracy: 0.9456 - val_loss: 1.2830 - val_accu-  
racy: 0.8076  
Epoch 30/35  
- 19s - loss: 0.2304 - accuracy: 0.9392 - val_loss: 1.4268 - val_accu-  
racy: 0.8068  
Epoch 31/35  
- 19s - loss: 0.1873 - accuracy: 0.9488 - val_loss: 1.3671 - val_accu-  
racy: 0.7943  
Epoch 32/35  
- 20s - loss: 0.1773 - accuracy: 0.9511 - val_loss: 1.1776 - val_accu-  
racy: 0.8311  
Epoch 33/35  
- 18s - loss: 0.1918 - accuracy: 0.9484 - val_loss: 1.3365 - val_accu-  
racy: 0.8197  
Epoch 34/35  
- 19s - loss: 0.2028 - accuracy: 0.9464 - val_loss: 1.2438 - val_accu-  
racy: 0.8233  
Epoch 35/35  
- 18s - loss: 0.1607 - accuracy: 0.9541 - val_loss: 1.1784 - val_accu-  
racy: 0.8260  
Baseline Error: 17.40%
```

Out[20]: <keras.engine.sequential.Sequential at 0x7f9808209650>

Conclusion: increasing units made the model worse. Maybe we should reduce?

```
In [23]: train_model(number_dense_layers=2, X_train= X_train, y_train=y_train,  
                    X_test=X_test, y_test= y_test, epochs=35, batch_size=100,  
                    units=750)
```

Train on 18620 samples, validate on 2552 samples

Epoch 1/35

- 5s - loss: 9.4356 - accuracy: 0.4048 - val_loss: 1.5659 - val_accuracy: 0.5157

Epoch 2/35

- 5s - loss: 1.1549 - accuracy: 0.6263 - val_loss: 1.2942 - val_accuracy: 0.6074

Epoch 3/35

- 5s - loss: 0.7875 - accuracy: 0.7426 - val_loss: 0.9384 - val_accuracy: 0.7288

Epoch 4/35

- 5s - loss: 0.5932 - accuracy: 0.8032 - val_loss: 0.8631 - val_accuracy: 0.7484

Epoch 5/35

- 5s - loss: 0.4716 - accuracy: 0.8456 - val_loss: 0.8662 - val_accuracy: 0.7594

Epoch 6/35

- 5s - loss: 0.3770 - accuracy: 0.8750 - val_loss: 0.9100 - val_accuracy: 0.7567

Epoch 7/35

- 5s - loss: 0.3466 - accuracy: 0.8868 - val_loss: 0.7817 - val_accuracy: 0.7978

Epoch 8/35

- 5s - loss: 0.2738 - accuracy: 0.9085 - val_loss: 0.8506 - val_accuracy: 0.7892

Epoch 9/35

- 5s - loss: 0.2450 - accuracy: 0.9175 - val_loss: 0.8997 - val_accuracy: 0.7900

Epoch 10/35

- 5s - loss: 0.2533 - accuracy: 0.9166 - val_loss: 0.8475 - val_accuracy: 0.8123

Epoch 11/35

- 5s - loss: 0.2845 - accuracy: 0.9076 - val_loss: 0.9785 - val_accuracy: 0.7908

Epoch 12/35

- 5s - loss: 0.3152 - accuracy: 0.9025 - val_loss: 0.9470 - val_accuracy: 0.8013

Epoch 13/35

- 5s - loss: 0.2323 - accuracy: 0.9245 - val_loss: 0.9522 - val_accuracy: 0.8100

Epoch 14/35

- 5s - loss: 0.2348 - accuracy: 0.9222 - val_loss: 0.9492 - val_accuracy: 0.8045

Epoch 15/35

- 5s - loss: 0.2034 - accuracy: 0.9351 - val_loss: 1.0116 - val_accuracy: 0.8154

Epoch 16/35

- 5s - loss: 0.2436 - accuracy: 0.9259 - val_loss: 1.0681 - val_accuracy: 0.8194

Epoch 17/35

- 5s - loss: 0.1927 - accuracy: 0.9403 - val_loss: 1.1365 - val_accuracy: 0.8021

Epoch 18/35

```
- 5s - loss: 0.2366 - accuracy: 0.9286 - val_loss: 1.0460 - val_accu-
racy: 0.8221
Epoch 19/35
- 5s - loss: 0.1395 - accuracy: 0.9531 - val_loss: 1.2403 - val_accu-
racy: 0.8088
Epoch 20/35
- 5s - loss: 0.1811 - accuracy: 0.9461 - val_loss: 1.1872 - val_accu-
racy: 0.8190
Epoch 21/35
- 5s - loss: 0.1582 - accuracy: 0.9501 - val_loss: 1.1546 - val_accu-
racy: 0.8111
Epoch 22/35
- 5s - loss: 0.1519 - accuracy: 0.9541 - val_loss: 1.1592 - val_accu-
racy: 0.8201
Epoch 23/35
- 5s - loss: 0.1774 - accuracy: 0.9475 - val_loss: 1.2742 - val_accu-
racy: 0.8009
Epoch 24/35
- 5s - loss: 0.2262 - accuracy: 0.9347 - val_loss: 1.2337 - val_accu-
racy: 0.8190
Epoch 25/35
- 5s - loss: 0.1826 - accuracy: 0.9450 - val_loss: 0.9523 - val_accu-
racy: 0.8350
Epoch 26/35
- 5s - loss: 0.1567 - accuracy: 0.9516 - val_loss: 1.0607 - val_accu-
racy: 0.8346
Epoch 27/35
- 5s - loss: 0.1481 - accuracy: 0.9549 - val_loss: 1.0552 - val_accu-
racy: 0.8342
Epoch 28/35
- 5s - loss: 0.1000 - accuracy: 0.9680 - val_loss: 0.9720 - val_accu-
racy: 0.8519
Epoch 29/35
- 5s - loss: 0.1531 - accuracy: 0.9555 - val_loss: 0.9742 - val_accu-
racy: 0.8260
Epoch 30/35
- 5s - loss: 0.1373 - accuracy: 0.9583 - val_loss: 1.1648 - val_accu-
racy: 0.8280
Epoch 31/35
- 5s - loss: 0.1387 - accuracy: 0.9597 - val_loss: 1.0769 - val_accu-
racy: 0.8221
Epoch 32/35
- 5s - loss: 0.1196 - accuracy: 0.9621 - val_loss: 1.0596 - val_accu-
racy: 0.8346
Epoch 33/35
- 5s - loss: 0.0944 - accuracy: 0.9699 - val_loss: 1.0122 - val_accu-
racy: 0.8476
Epoch 34/35
- 5s - loss: 0.1280 - accuracy: 0.9642 - val_loss: 1.3175 - val_accu-
racy: 0.8107
Epoch 35/35
- 5s - loss: 0.1536 - accuracy: 0.9572 - val_loss: 0.9499 - val_accu-
racy: 0.8378
Baseline Error: 16.22%
```

Out[23]: <keras.engine.sequential.Sequential at 0x7f98097ef0d0>

```
The best model remains the one with 2 layers and batch size 100  
Could adding more layers help?
```

```
In [24]: train_model(number_dense_layers=3, X_train= X_train, y_train=y_train,  
                    X_test=X_test, y_test= y_test, epochs=35, batch_size=100,  
                    units=2000)
```

Train on 18620 samples, validate on 2552 samples

Epoch 1/35

- 24s - loss: 76.6539 - accuracy: 0.4060 - val_loss: 2.0274 - val_accuracy: 0.5196

Epoch 2/35

- 23s - loss: 1.1814 - accuracy: 0.6670 - val_loss: 1.4257 - val_accuracy: 0.6399

Epoch 3/35

- 26s - loss: 0.7189 - accuracy: 0.7777 - val_loss: 1.2475 - val_accuracy: 0.6963

Epoch 4/35

- 26s - loss: 0.4779 - accuracy: 0.8485 - val_loss: 1.1640 - val_accuracy: 0.7183

Epoch 5/35

- 25s - loss: 0.3239 - accuracy: 0.8917 - val_loss: 1.1322 - val_accuracy: 0.7426

Epoch 6/35

- 25s - loss: 0.2731 - accuracy: 0.9070 - val_loss: 1.1351 - val_accuracy: 0.7586

Epoch 7/35

- 23s - loss: 0.2297 - accuracy: 0.9241 - val_loss: 1.2070 - val_accuracy: 0.7539

Epoch 8/35

- 23s - loss: 0.2151 - accuracy: 0.9277 - val_loss: 1.1571 - val_accuracy: 0.7798

Epoch 9/35

- 26s - loss: 0.2649 - accuracy: 0.9170 - val_loss: 1.2437 - val_accuracy: 0.7637

Epoch 10/35

- 25s - loss: 0.3291 - accuracy: 0.9045 - val_loss: 1.1487 - val_accuracy: 0.7763

Epoch 11/35

- 25s - loss: 0.3129 - accuracy: 0.9062 - val_loss: 1.3616 - val_accuracy: 0.7731

Epoch 12/35

- 26s - loss: 0.2545 - accuracy: 0.9230 - val_loss: 1.4635 - val_accuracy: 0.7672

Epoch 13/35

- 24s - loss: 0.3164 - accuracy: 0.9131 - val_loss: 1.1724 - val_accuracy: 0.7861

Epoch 14/35

- 25s - loss: 0.2145 - accuracy: 0.9367 - val_loss: 1.0639 - val_accuracy: 0.8201

Epoch 15/35

- 26s - loss: 0.2471 - accuracy: 0.9308 - val_loss: 1.1629 - val_accuracy: 0.8123

Epoch 16/35

- 25s - loss: 0.2221 - accuracy: 0.9373 - val_loss: 1.3010 - val_accuracy: 0.8041

Epoch 17/35

- 25s - loss: 0.3314 - accuracy: 0.9134 - val_loss: 1.4322 - val_accuracy: 0.7778

Epoch 18/35

```
- 25s - loss: 0.2783 - accuracy: 0.9234 - val_loss: 1.3126 - val_accu-  
racy: 0.7876  
Epoch 19/35  
- 26s - loss: 0.2070 - accuracy: 0.9412 - val_loss: 1.5889 - val_accu-  
racy: 0.7810  
Epoch 20/35  
- 27s - loss: 0.2516 - accuracy: 0.9318 - val_loss: 1.1914 - val_accu-  
racy: 0.8111  
Epoch 21/35  
- 26s - loss: 0.2074 - accuracy: 0.9424 - val_loss: 1.0575 - val_accu-  
racy: 0.8252  
Epoch 22/35  
- 28s - loss: 0.1938 - accuracy: 0.9466 - val_loss: 1.3285 - val_accu-  
racy: 0.7990  
Epoch 23/35  
- 26s - loss: 0.2396 - accuracy: 0.9359 - val_loss: 1.2530 - val_accu-  
racy: 0.8147  
Epoch 24/35  
- 24s - loss: 0.2447 - accuracy: 0.9334 - val_loss: 1.3976 - val_accu-  
racy: 0.7923  
Epoch 25/35  
- 23s - loss: 0.1591 - accuracy: 0.9532 - val_loss: 1.1413 - val_accu-  
racy: 0.8197  
Epoch 26/35  
- 23s - loss: 0.1945 - accuracy: 0.9471 - val_loss: 1.1266 - val_accu-  
racy: 0.8186  
Epoch 27/35  
- 24s - loss: 0.1773 - accuracy: 0.9492 - val_loss: 1.1547 - val_accu-  
racy: 0.8111  
Epoch 28/35  
- 24s - loss: 0.1773 - accuracy: 0.9503 - val_loss: 1.0609 - val_accu-  
racy: 0.8245  
Epoch 29/35  
- 23s - loss: 0.1776 - accuracy: 0.9528 - val_loss: 1.1838 - val_accu-  
racy: 0.8241  
Epoch 30/35  
- 23s - loss: 0.1860 - accuracy: 0.9483 - val_loss: 1.1191 - val_accu-  
racy: 0.8229  
Epoch 31/35  
- 23s - loss: 0.1738 - accuracy: 0.9507 - val_loss: 1.0110 - val_accu-  
racy: 0.8248  
Epoch 32/35  
- 23s - loss: 0.2311 - accuracy: 0.9374 - val_loss: 1.1553 - val_accu-  
racy: 0.8237  
Epoch 33/35  
- 23s - loss: 0.1761 - accuracy: 0.9505 - val_loss: 1.2413 - val_accu-  
racy: 0.8241  
Epoch 34/35  
- 24s - loss: 0.2208 - accuracy: 0.9416 - val_loss: 0.9756 - val_accu-  
racy: 0.8346  
Epoch 35/35  
- 23s - loss: 0.2263 - accuracy: 0.9402 - val_loss: 1.0766 - val_accu-  
racy: 0.8225  
Baseline Error: 17.75%
```

Out[24]: <keras.engine.sequential.Sequential at 0x7f980bb39c90>

Speech conclusion

- The best speech model was a 2-layer model with 35 epochs. We probably need to introduce a type of early stopping where the number of epochs is halted after observing one or two decreases in validation accuracy.
- Although 1000 unit seems arbitray, neither decreasing, nor increase the number seemed to produce a more accurate model

Natural Language Processing: IMDb review sentiment analysis

```
In [8]: from db_utils import get_imdb_dataset
(X_train, y_train), (X_test, y_test) = get_imdb_dataset()
```

```
In [88]: train_model(number_dense_layers=1, X_train= X_train, y_train=y_train,
X_test=X_test, y_test= y_test, epochs=10, batch_size=200,
units=1000)
```

Train on 25000 samples, validate on 25000 samples

Epoch 1/10

- 239s - loss: 0.3478 - accuracy: 0.8502 - val_loss: 0.2992 - val_accuracy: 0.8777

Epoch 2/10

- 285s - loss: 0.0556 - accuracy: 0.9855 - val_loss: 0.3229 - val_accuracy: 0.8769

Epoch 3/10

- 235s - loss: 0.0129 - accuracy: 0.9986 - val_loss: 0.3615 - val_accuracy: 0.8776

Epoch 4/10

- 246s - loss: 0.0044 - accuracy: 0.9999 - val_loss: 0.3898 - val_accuracy: 0.8775

Epoch 5/10

- 256s - loss: 0.0024 - accuracy: 1.0000 - val_loss: 0.4160 - val_accuracy: 0.8778

Epoch 6/10

- 277s - loss: 0.0017 - accuracy: 1.0000 - val_loss: 0.4308 - val_accuracy: 0.8780

Epoch 7/10

- 292s - loss: 8.7237e-04 - accuracy: 1.0000 - val_loss: 0.4467 - val_accuracy: 0.8773

Epoch 8/10

- 292s - loss: 6.2179e-04 - accuracy: 1.0000 - val_loss: 0.4601 - val_accuracy: 0.8772

Epoch 9/10

- 239s - loss: 4.6525e-04 - accuracy: 1.0000 - val_loss: 0.4723 - val_accuracy: 0.8767

Epoch 10/10

- 246s - loss: 3.5998e-04 - accuracy: 1.0000 - val_loss: 0.4832 - val_accuracy: 0.8764

Baseline Error: 12.36%

```
Out[88]: <keras.engine.sequential.Sequential at 0x7f859df10410>
```

Conclusions

- Accuracy on the training set is 100% while accuracy on the validation set is only 87.7%
- It seems the model is overfitting. As such, longer training with an identical model doesn't seem reasonable.


```
In [9]: # adding an additional dense layer
train_model(number_dense_layers=2, X_train= X_train, y_train=y_train,
            X_test=X_test, y_test= y_test, epochs=10, batch_size=200,
            units=1000)
```

Train on 25000 samples, validate on 25000 samples

Epoch 1/10

- 199s - loss: 0.4523 - accuracy: 0.8070 - val_loss: 0.2909 - val_accuracy: 0.8794

Epoch 2/10

- 193s - loss: 0.0651 - accuracy: 0.9795 - val_loss: 0.3443 - val_accuracy: 0.8681

Epoch 3/10

- 189s - loss: 0.0039 - accuracy: 0.9996 - val_loss: 0.4677 - val_accuracy: 0.8736

Epoch 4/10

- 191s - loss: 3.6917e-04 - accuracy: 1.0000 - val_loss: 0.5322 - val_accuracy: 0.8720

Epoch 5/10

- 187s - loss: 7.5469e-04 - accuracy: 1.0000 - val_loss: 0.5352 - val_accuracy: 0.8756

Epoch 6/10

- 178s - loss: 8.8936e-05 - accuracy: 1.0000 - val_loss: 0.5627 - val_accuracy: 0.8756

Epoch 7/10

- 182s - loss: 5.6146e-05 - accuracy: 1.0000 - val_loss: 0.5849 - val_accuracy: 0.8756

Epoch 8/10

```
-----
--
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-9-4c52fc23b648> in <module>
      2 train_model(number_dense_layers=2, X_train= X_train, y_train=y_train,
      3 X_test=X_test, y_test= y_test, epochs=10, batch_size=200,
----> 4 units=1000)

<ipython-input-2-536cc44a0f99> in train_model(number_dense_layers, X_train, y_train, X_test, y_test, epochs, batch_size, units)
     22             y_train.shape[1], units)
     23     model.fit(X_train, y_train, validation_data=(X_test, y_test),
--> 24                 epochs=epochs, batch_size=batch_size, verbose=2)
     25     scores = model.evaluate(X_test, y_test, verbose=2)
     26     print("Baseline Error: %.2f%%" % (100-scores[1]*100))

/opt/anaconda3/envs/deep37/lib/python3.7/site-packages/keras/engine/training.py in fit(self, x, y, batch_size, epochs, verbose, callbacks, validation_split, validation_data, shuffle, class_weight, sample_weight, initial_epoch, steps_per_epoch, validation_steps, validation_freq, max_queue_size, workers, use_multiprocessing, **kwargs)
    1237         steps_per_epoch=steps_per
    _epoch,
    1238         validation_steps=validati
on_steps,
-> 1239         validation_freq=validatio
```

```

n_freq)
1240
1241     def evaluate(self,

/opt/anaconda3/envs/deep37/lib/python3.7/site-packages/keras/engine/train
ing_arrays.py in fit_loop(model, fit_function, fit_inputs, out_labels, ba
tch_size, epochs, verbose, callbacks, val_function, val_inputs, shuffle,
    initial_epoch, steps_per_epoch, validation_steps, validation_freq)
194         ins_batch[i] = ins_batch[i].toarray()
195
--> 196         outs = fit_function(ins_batch)
197         outs = to_list(outs)
198         for l, o in zip(out_labels, outs):

/opt/anaconda3/envs/deep37/lib/python3.7/site-packages/tensorflow_core/py
thon/keras/backend.py in __call__(self, inputs)
3738         value = math_ops.cast(value, tensor.dtype)
3739         converted_inputs.append(value)
-> 3740     outputs = self._graph_fn(*converted_inputs)
3741
3742     # EagerTensor.numpy() will often make a copy to ensure memory
safety.

/opt/anaconda3/envs/deep37/lib/python3.7/site-packages/tensorflow_core/py
thon/eager/function.py in __call__(self, *args, **kwargs)
1079     TypeError: For invalid positional/keyword argument combinat
ions.
1080     """
-> 1081     return self._call_impl(args, kwargs)
1082
1083     def _call_impl(self, args, kwargs, cancellation_manager=None):

/opt/anaconda3/envs/deep37/lib/python3.7/site-packages/tensorflow_core/py
thon/eager/function.py in _call_impl(self, args, kwargs, cancellation_man
ager)
1119         raise TypeError("Keyword arguments {} unknown. Expected
{}".format(
1120             list(kwargs.keys()), list(self._arg_keywords)))
-> 1121     return self._call_flat(args, self.captured_inputs, cancellati
on_manager)
1122
1123     def _filtered_call(self, args, kwargs):

/opt/anaconda3/envs/deep37/lib/python3.7/site-packages/tensorflow_core/py
thon/eager/function.py in _call_flat(self, args, captured_inputs, cancell
ation_manager)
1222         if executing_eagerly:
1223             flat_outputs = forward_function.call(
-> 1224                 ctx, args, cancellation_manager=cancellation_manager)
1225         else:
1226             gradient_name = self._delayed_rewrite_functions.register()

/opt/anaconda3/envs/deep37/lib/python3.7/site-packages/tensorflow_core/py
thon/eager/function.py in call(self, ctx, args, cancellation_manager)
509         inputs=args,
510         attrs=("executor_type", executor_type, "config_prot
o", config),

```

```

--> 511             ctx=ctx)
      512         else:
      513             outputs = execute.execute_with_cancellation(

/opt/anaconda3/envs/deep37/lib/python3.7/site-packages/tensorflow_core/py
thon/eager/execute.py in quick_execute(op_name, num_outputs, inputs, attr
s, ctx, name)
      59         tensors = pywrap_tensorflow.TFE_Py_Execute(ctx._handle, devic
e_name,
      60                                                         op_name, inputs, a
ttrs,
--> 61                                                         num_outputs)
      62     except core._NotOkStatusException as e:
      63         if name is not None:

```

KeyboardInterrupt:

Conclusions on Natural Language Processing

- Accuracy on the training set is 100% while accuracy on the validation set is only 87.7%
- It seems the model is overfitting. As such, longer training with an identical model doesn't seem reasonable.
- We therefore halted training.

Overall conclusions

- We ran an identical set of functions on a vision, speech and NLP model.
- The fact that a couple dozen lines of code can produce a model across a variety of use cases is pretty amazing
- Optimizing hyperparameters seems trial and error so far. Increasing the number of layers did improve speech but not vision. Increasing epochs seemed to have a positive effect on both vision and speech. But NLP seems to need a different model spec in order to improve performance.
- Training time can go up significantly, even with these small datasets. We tried the vision model on google collab, and training times decreased materially.